


User Manual.

Six agents.

AI-powered content generation for multi-channel brands.
One idea. Every channel. Six AI agents. Forever.



CHAPTER ONE

Welcome to GEM

GEM is a self-hosted AI content engine that turns a single content idea into a complete, platform-ready package for every channel you manage. It runs in any modern browser — entirely local on Ollama, or against any cloud LLM (OpenAI, Anthropic, OpenRouter) with your own key. No GEM servers. No data leaves your computer unless you choose a cloud provider.

Where most AI tools give you a generic blog post or a single caption, GEM orchestrates six specialist agents in sequence to produce: brand strategy, viral topic intelligence, a structured framework, a humanized voiceover script, master image and video prompts for every aspect ratio, and individually tailored copy for each of your YouTube, Facebook, Instagram, TikTok, LinkedIn, and blog channels.

What makes GEM different

	Generic AI Tools	GEM Content Engine
Cost	\$20–\$100 / month forever	One-time purchase · \$9 Founder's Edition
Privacy	Your data trains their AI	100% local via Ollama (or BYO cloud key)
Output	One generic blog post	6 agents · 17+ asset types per run
Channel-aware	Same copy for everywhere	Adapted per channel and per audience
Image prompts	Basic text descriptions	Master prompts for MJ v6, DALL-E 3, FLUX.1
Video prompts	Not included	Kling AI + RunwayML, scene-by-scene
Script quality	Robotic, listicle-style	Humanized, conversational, graded by Quality Bar
Trending data	LLM guesses	Live Reddit / HN (with r/popular fallback for Google Trends)

CHAPTER TWO

The Six AI Agents

GEM's pipeline runs six specialist agents in sequence. Each agent has one job, hands its structured output down the line, and produces a downloadable report. The Quality Bar grades each agent's output against six rules before the next agent runs.

1. Brand Brain — Channel strategy · voice · pillars

Analyzes your channels and niche, produces brand voice doc, 3–5 content pillars, audience definition, competitive positioning, and the anchor hook used by every downstream agent.

2. Trend Scout — Viral topic discovery

Pulls live signals from Reddit and HackerNews (CORS-friendly). Google Trends falls back to r/popular (documented in Chapter 7). Ranks 5–6 topic candidates by viral score, estimated views, competition, and platform fit.

3. Content Architect — Structure & retention hooks

Designs section-by-section structure with timestamped scenes, retention hooks at 30s / 2m / 5m, SEO title, thumbnail concept, upload timing, and revenue estimate per 1,000 views.

4. Script Writer — Humanized voiceover

Writes scripts that pass the Quality Bar two-second-hook test: natural speech, contractions, [pause] markers, rhetorical questions, sentences starting with "And/But/Look", tension-before-reveal pacing.

5. Visual Director — Master prompts (the flagship agent)

Ultra-detailed 4K master prompts for Midjourney v6, DALL-E 3 / Imagen, and FLUX.1 across every format you need: thumbnail, IG post, IG story, FB cover, YT banner, blog hero. Plus scene-by-scene Kling AI and RunwayML video prompts.

6. Channel Adapter — Per-channel content

Takes the master content and produces individually adapted copy for every channel you have configured. Each Facebook page gets its own caption. Each Instagram page gets its own Reel hook and hashtags. YouTube SEO package, blog intro — all platform-perfect.

CHAPTER THREE

Setup & Installation

System requirements

Component	Minimum	Recommended
Operating system	Win 10 / macOS 11 / Linux	Any 64-bit OS
RAM (Ollama mode)	8 GB (7B models)	16–32 GB (14B+ models)
RAM (cloud mode)	4 GB	8 GB
Disk	10 GB free	50 GB for multiple Ollama models
CPU	Modern x86_64 / ARM64	Apple Silicon / Ryzen / Core i7+
GPU	Not required	NVIDIA GPU dramatically speeds inference
Browser	Chrome 100+ / Firefox 100+	Latest Chrome or Edge

Local mode (Ollama) — 5 steps

- ❖ **Step 1.** Install Ollama from ollama.ai. On macOS / Linux: `curl -fsSL https://ollama.ai/install.sh | sh`
- ❖ **Step 2.** Pull a model. For 16 GB+ systems: `ollama pull qwen2.5:14b`. For 8 GB systems: `ollama pull qwen2.5:7b`.
- ❖ **Step 3.** Start Ollama with CORS so the GEM browser interface can talk to it: `OLLAMA_ORIGINS="*" ollama serve`
- ❖ **Step 4.** Open `gem.html` in your browser. Click Configure → Ollama → paste `http://localhost:11434` → choose your model → Test connection. Green dot means success.
- ❖ **Step 5.** Add your channels in the left sidebar. Channels live in your browser only — nothing is sent anywhere.

Cloud mode (BYO key) — 1 step

Open `gem.html` → Settings → Provider → pick OpenAI, Anthropic, or OpenRouter → paste your key. A full pipeline run costs roughly \$0.02–\$0.05 on gpt-4o-mini or \$0.12 on Claude Sonnet. Keys live in your browser only; GEM never sees them.

CHAPTER FOUR

Daily Workflow

Once setup is complete, your daily workflow takes 5–10 minutes for a complete content drop. Quality Bar checks run automatically between agents — you can't ship a generic hook by accident.

- **1. Pick a niche.** On the welcome screen, choose your niche from the 12 templates or type a custom one. The more specific the input, the sharper the output.
- **2. Toggle Trending Mode (optional).** On the welcome screen, flip the Trending Mode toggle to pull live Reddit / HackerNews signals. Google Trends falls back to r/popular per the disclosure in Chapter 7.
- **3. Hit Generate.** Click Generate Content Drop. Watch the agent activity log on the right — each agent reports in as it completes. Quality Bar gates are visible between agents.
- **4. Review the drop.** You'll get script + image prompts + video prompts + per-channel captions + carousels + a 30-day calendar slot suggestion. Approve or edit each asset.
- **5. Export.** Designer-grade PDF, JSON (.gem.json), or copy-paste directly into your scheduler. The PDF is brand-styled.

CHAPTER FIVE

Complete Feature Reference

Each pipeline run produces 17+ asset types. Below is the full list by agent.

Brand Brain

- ✓ Brand voice document
- ✓ Content pillars (3–5)
- ✓ Audience definition + persona
- ✓ Competitive positioning
- ✓ Anchor hook for the pipeline

Trend Scout

- ✓ Ranked topic candidates (5–6)
- ✓ Viral score + estimated views
- ✓ Live Reddit / HackerNews signals
- ✓ Google Trends via r/popular fallback

Content Architect

- ✓ Section-by-section structure
- ✓ Timestamped retention hooks
- ✓ SEO title + thumbnail concept
- ✓ Optimal upload timing

Script Writer

- ✓ Humanized voiceover script
- ✓ Hook → setup → reveal → close pacing
- ✓ Two-second-hook check

Visual Director

- ✓ Midjourney v6 master prompts (every aspect ratio)
- ✓ DALL-E 3 / Imagen prompts (5 per piece)
- ✓ FLUX.1 prompts
- ✓ Kling AI scene prompts
- ✓ RunwayML scene prompts

- ✓ Multi-slide Instagram + LinkedIn carousels

Channel Adapter

- ✓ YouTube SEO title + description + tags
- ✓ Instagram caption + Reel hook + cover text
- ✓ Facebook hook + caption
- ✓ TikTok / Shorts hook
- ✓ LinkedIn post
- ✓ Blog title + meta + intro

CHAPTER SIX

Power Features (v1.7)

Quality Bar (new in v1.7)

Six rules every drop must clear: one clear idea per post, two-second hook, specific examples not generic advice, on-brand visuals, business-goal CTA, human approval beat. Each agent's output is graded automatically before the next runs.

"Today's Brief" sample (new in v1.7)

The live landing page now ships a real "AI agents replacing office work" pack so prospects see exactly what GEM hands them. You can render the same demo locally by loading the bundled .gem.json from the templates folder.

Roadmap voting (new in v1.7)

Founder's Edition buyers vote on what ships next. The four big modules in development: daily auto-brief, approval workflow, direct publish to Buffer / Later / Meta, performance loop.

Trending Mode

Live scraping of Reddit and HackerNews directly in your browser. Google Trends has stricter CORS and falls back to r/popular as a trending proxy — covers 95% of niches. Run a one-line local proxy for unrestricted Google Trends; see Chapter 7.

Photo Post Prompts

5 ready-to-paste ChatGPT / DALL-E 3 / Imagen prompts per piece for IG Feed, IG Reel, FB, TikTok, LinkedIn.

Carousel Generator

Multi-slide IG + LinkedIn carousels with role (Hook / Setup / Reveal / CTA), slide headline, and ChatGPT image prompt per slide.

Brand Kit

Color picker with hex input, native picker, and 8 preset palettes (Forest, Ocean, Sunset, Royal, Rose, Mono, Earth, Mint). Brand flows into every Visual prompt.

My Library

Auto-saved content history. Click any past piece to instantly reload it. Search and filter coming in v1.8.

30-Day Calendar

AI-generated content calendar across all your platforms in 60 seconds. Drag-to-reorder, with per-platform best-time-to-post suggestions.

Tone Presets

Balanced / Educational / Entertaining / Controversial / Inspirational. One-click voice swap.

Channel Export / Import

JSON download / upload of channel configs. Cross-device sync, team handoff, backup.

CHAPTER SEVEN

Troubleshooting

Ollama "Connection failed" with green-dot test

You probably forgot to start Ollama with CORS enabled. Quit Ollama and relaunch with `OLLAMA_ORIGINS="*" ollama serve` in a terminal.

Trending Mode shows only r/popular signals

This is expected behavior when Google Trends is blocked by CORS — documented in the public FAQ. Reddit and HackerNews still work directly. For unrestricted Google Trends, run a one-line local proxy: `npx cors-anywhere` on port 8080, then point Trending Mode at `http://localhost:8080`.

Cloud mode shows "Invalid API key"

Check the key starts with the right prefix (`sk-` for OpenAI, `sk-ant-` for Anthropic). For OpenRouter, your key starts with `sk-or-`. Keys are stored in your browser only — paste fresh ones if you rotate.

Pipeline stalls between agents

Quality Bar may be rejecting an agent's output. Open the agent activity log on the right — it will show which rule failed and let you regenerate that specific agent without re-running upstream.

Image prompts come out generic

Your Brand Kit is probably empty. Fill in colors + style notes + 2–3 reference adjectives before running the pipeline; the Visual agent uses them on every prompt.

PDF export is blank

Browser print blocked custom fonts. Switch to Chrome / Edge and re-export. If still blank, open Settings → Export → "Use system fonts" and try again.

Saved .gem.json won't load

File was edited externally. Re-export from the Library with the original file selected, or restore from My Library auto-save (last 30 runs are kept).

CHAPTER EIGHT

License & Final Notes

Your license

Your purchase of GEM grants you a lifetime license to use the tool for personal and commercial content creation. Use it for your own brand, your clients, your agency, and your team. The single HTML file is yours to keep, modify, and integrate into your workflow forever — no recurring fees, no expiration, no telemetry. Founder's Edition includes the commercial-use license.

What's included

- ✓ The complete gem.html file (single self-contained app)
- ✓ This v1.7 user manual PDF
- ✓ Lifetime updates within the v1.x major version
- ✓ Discord community access for users
- ✓ Direct line to the founder for feedback (Founder's Edition)
- ✓ Email support for setup issues (first 30 days)

Privacy promise

GEM runs 100% locally on Ollama mode. Your channel configurations, content ideas, generated scripts, and prompts never leave your computer. We have no servers, no analytics, no telemetry. The only network calls made by GEM in Ollama mode are between your browser and your local Ollama instance. In cloud mode, prompts go to your chosen provider (OpenAI / Anthropic / OpenRouter) and nowhere else — GEM has no intermediary server.

A note from the creator

GEM exists because the AI tooling market is structured against creators. You pay \$20–\$100 a month for tools that own your data, lock you into their cloud, and produce output indistinguishable from a thousand other creators using the same prompts. GEM flips that: one fair price, your data stays yours, and the output is uniquely tuned to your channels. Make great content. Build your audience. Keep your data. — The GEM team.